Independence and Bayesian Networks (Part 1)

Wenhu Chen

Lecture 3

Outline

Learning Goals

Unconditional and Conditional Independence

Examples of Bayesian Networks

Why Bayesian Networks

Representing the Joint Distribution

Independence in Three Key Structures

Revisiting Learning Goals

Learning Goals

- Given a probabilistic model, determine if two variables are unconditionally independent, or conditionally independent given a third variable.
- Give examples of deriving a compact representation of a joint distribution by using independence assumptions.
- Describe components of a Bayesian network.
- Compute a joint probability given a Bayesian network.
- Explain the independence relationships in the three key structures.

CS 486/686: Intro to AI Lecturer: Wenhu Chen 3 / 53

Learning Goals

Unconditional and Conditional Independence

Examples of Bayesian Networks

Why Bayesian Networks

Representing the Joint Distribution

Independence in Three Key Structures

Revisiting Learning Goals

Definition ((unconditional) independence)

X and Y are (unconditionally) independent iff

$$P(X|Y) = P(X)$$

$$P(Y|X) = P(Y)$$

$$P(X \land Y) = P(X)P(Y)$$

Learning Y does NOT influence your belief about X.

Definition ((unconditional) independence)

X and Y are (unconditionally) independent iff

$$P(X|Y) = P(X)$$

$$P(Y|X) = P(Y)$$

$$P(X \land Y) = P(X)P(Y)$$

Learning Y does NOT influence your belief about X.

 \rightarrow Convert between the two equations.

To specify joint probability, it is sufficient to specify the individual probabilities.

CS 486/686: Intro to AI

Definition ((unconditional) independence)

X and Y are (unconditionally) independent iff

$$P(X|Y) = P(X)$$

$$P(Y|X) = P(Y)$$

$$P(X \land Y) = P(X)P(Y)$$

Learning Y does NOT influence your belief about X.

Definition ((unconditional) independence)

X and Y are (unconditionally) independent iff

$$P(X|Y) = P(X)$$

$$P(Y|X) = P(Y)$$

$$P(X \land Y) = P(X)P(Y)$$

Learning Y does NOT influence your belief about X.

 \rightarrow To justify that

$$P(X \wedge Y) = P(X)P(Y)$$

we need to make four comparisons.

CS 486/686: Intro to AI

Lecturer: Wenhu Chen

Definition (conditional independence)

 \boldsymbol{X} and \boldsymbol{Y} are conditionally independent given \boldsymbol{Z} if

$$P(X|Y \wedge Z) = P(X|Z).$$

$$P(Y|X \wedge Z) = P(Y|Z).$$

$$P(Y \wedge X|Z) = P(Y|Z)P(X|Z).$$

Learning Y does NOT influence your belief about X if you already know Z.

CS 486/686: Intro to AI

Definition (conditional independence)

 \boldsymbol{X} and \boldsymbol{Y} are conditionally independent given \boldsymbol{Z} if

$$P(X|Y \wedge Z) = P(X|Z).$$

$$P(Y|X \wedge Z) = P(Y|Z).$$

$$P(Y \wedge X|Z) = P(Y|Z)P(X|Z).$$

Learning Y does NOT influence your belief about X if you already know Z.

 $\to X$ is conditionally independent of Y given Z.

Independence does not imply conditional independence, and vice versa.

CS 486/686: Intro to AI Lecturer: Wenhu Chen 7 / 53

Definition (conditional independence)

X and Y are conditionally independent given Z if

$$P(X|Y \wedge Z) = P(X|Z).$$

$$P(Y|X \wedge Z) = P(Y|Z).$$

$$P(Y \wedge X|Z) = P(Y|Z)P(X|Z).$$

Learning Y does NOT influence your belief about X if you already know Z.

Definition (conditional independence)

X and Y are conditionally independent given Z if

$$P(X|Y \wedge Z) = P(X|Z).$$

$$P(Y|X \wedge Z) = P(Y|Z).$$

$$P(Y \wedge X|Z) = P(Y|Z)P(X|Z).$$

Learning Y does NOT influence your belief about X if you already know Z.

 \rightarrow To justify that

$$P(X \land Y|Z) = P(X|Z)P(Y|Z)$$

we need to make eight comparisons.

CS 486/686: Intro to AI

Q: Consider a model with three random variables, A, B, C. What is the minimum number of probabilities required to specify the joint distribution?

- (A) 3
- (B) 7
- (C) 8
- (D) 16

Q: Consider a model with three random variables, A, B, C. What is the minimum number of probabilities required to specify the joint distribution?

- (A) 3
- (B) 7
- (C) 8
- (D) 16
- \rightarrow (C) $P(A), P(B|A), P(C|A \land B)$. 1 + 2 + 4 = 7 probabilities

Draw a graph to prove it to yourself.

Q: Consider a model with three random variables, A, B, C. What is the minimum number of probabilities required to specify the joint distribution?

- (A) 3
- (B) 7
- (C) 8
- (D) 16
- \rightarrow (C) $P(A, B, C), P(\neg A, B, C), \cdots, P(\neg A, \neg B, \neg C)$. A total of 8 - 1 = 7 probabilities

Q: Consider a model with three random variables, A, B, C. Assume that A, B, and C are independent. What is the minimum number of probabilities required to specify the joint distribution?

- (A) 3
- (B) 7
- (C) 8
- (D) 16

Q: Consider a model with three random variables, A, B, C. Assume that A, B, and C are independent. What is the minimum number of probabilities required to specify the joint distribution?

- (A) 3
- (B) 7
- (C) 8
- (D) 16
- \rightarrow (A) P(A), P(B), P(C), 1 + 1 + 1 = 3 probabilities

Draw a graph to prove it to yourself.

Q: Consider a model with three boolean random variables, A,B,C. Assume that A and B are conditionally independent given C. What is the minimum number of probabilities required to specify the joint distribution?

- (A) 4
- (B) 5
- (C) 7
- (D) 11

Q: Consider a model with three boolean random variables, A,B,C. Assume that A and B are conditionally independent given C. What is the minimum number of probabilities required to specify the joint distribution?

- (A) 4
- (B) 5
- (C) 7
- (D) 11
- → (B) P(C), P(A|C), P(B|C). 1 + 2 + 2 = 5 probabilities

Draw a graph to prove it to yourself.

Q: Given the joint probability distribution about A, B, and C, how many comparisons do we need to make to justify A and B are independent given C?

- (A) 1
- (B) 4
- (C) 8
- (D) 6

Q: Given the joint probability distribution about A, B, and C, how many comparisons do we need to make to justify A and B are independent given C?

- (A) 1
- (B) 4
- (C) 8
- (D) 6

$$\rightarrow$$
 (C) $p(B,A|C) = p(B|C) * p(A|C)$ and $p(B,\neg A|C) = p(B|C) * p(\neg B|C)$

... A total of 8 equalities!

Is this true?

Q: Given the joint probability distribution about A, B, and C, how many comparisons do we need to make to justify A and B are independent given C?

- (A) 1
- (B) 4
- (C) 8
- (D) 6

Q: Given the joint probability distribution about A, B, and C, how many comparisons do we need to make to justify A and B are independent given C?

- (A) 1
- (B) 4
- (C) 8
- (D) 6
- \rightarrow If we already know

$$\begin{array}{l} p(B,A|C)=p(B|C)*p(A|C); p(\neg B,A|C)=\\ p(\neg B|C)*p(A|C); p(B,\neg A|C)=p(B|C)*p(\neg A|C). \text{ Do we still need to compare }p(\neg B,\neg A|C) \text{ and }p(\neg B|C)*p(\neg A|C)? \end{array}$$

Probably not, the answer is (D).

Q: Read the table to understand whether B and C are independent given A.

Α	В	С	Prob
Т	Т	Т	0.16
Τ	Т	F	0.16
Τ	F	Т	0.24
Τ	F	F	0.24
F	Т	Т	0.012
F	Т	F	0.008
F	F	Т	0.108
F	F	F	0.072

- (A) B and C are independent given A
- (B) B and C are not independent given A

Read the table to understand whether B and C are independent given A.

Α	В	C	Prob
Т	Т	Т	0.16
Τ	Т	F	0.16
Т	F	Т	0.24
Т	F	F	0.24
F	Т	Т	0.012
F	Т	F	0.008
F	F	Т	0.108
F	F	F	0.072

- ightharpoonup Compute p(B, C|A)
- ▶ Compute p(B|A) and p(C|A)
- $\blacktriangleright \text{ Verify } p(B,C|A) = p(B|A) * p(C|A)$

Q #3b: Step-by-Step Derivation p(B, C|A)

Α	В	С	Prob
Т	Т	Т	0.16
Т	Т	F	0.16
Т	F	Т	0.24
Т	F	F	0.24
F	Т	Т	0.012
F	Т	F	0.008
F	F	Т	0.108
F	F	F	0.072

Table: Merging p(A, B, C).

- p(B, C|A) = p(B, C, A)/p(A)
- p(A) = (0.16 + 0.16 + 0.24 + 0.24, 0.012 + 0.008 + 0.108 + 0.072) = (0.8, 0.2)

Q #3b: Step-by-Step Derivation p(B, C|A)

В	С	(A)	Prob
Т	Т	Т	0.16 / 0.8 = 0.2
Т	F	Т	0.16 / 0.8 = 0.2
F	Т	Т	0.24 / 0.8 = 0.3
F	F	Τ	0.24 / 0.8 = 0.3
Т	Т	F	0.012 / 0.2 = 0.06
Т	F	F	0.008 / 0.2 = 0.04
F	Т	F	0.108 / 0.2 = 0.54
F	F	F	0.072 / 0.2 = 0.36

Table: Computing p(B, C|A).

- p(B, C|A) = p(B, C, A)/p(A)
- p(A) = (0.8, 0.2)
- ightharpoonup p(B,C|A) is displayed in the table

Q #3b: Step-by-Step Derivation

Α	В	С	Prob
Т	Т	Т	0.16
Т	Т	F	0.16
Т	F	Т	0.24
Т	F	F	0.24
F	Т	Т	0.012
F	Т	F	0.008
F	F	Т	0.108
F	F	F	0.072

Table: Merge p(A, B, C)

ightharpoonup Marginalizing over variable C

Q #3b: Step-by-Step Derivation p(B|A)

Α	В	Prob
Т	Т	0.32
Т	F	0.48
F	Т	0.02
F	F	0.18

Table: Computing p(A, B)

- ightharpoonup Marginalizing over variable C
- ▶ Joint p(A, B) is displayed in the table

Q #3b: Step-by-Step Derivation p(B|A)

В	(A)	Prob
Т	Т	0.32 / 0.8 = 0.4
F	Т	0.48 / 0.8 = 0.6
Т	F	0.02 / 0.2 = 0.1
F	F	0.18 / 0.2 = 0.9

Table: Computing p(B|A)

- ightharpoonup Marginalizing over variable C
- lacktriangle Conditional p(B|A) is displayed in the table

Q #3b: Step-by-Step Derivation p(C|A)

Α	В	С	Prob
Т	Т	Т	0.16
Т	F	Т	0.24
Т	Т	F	0.16
Т	F	F	0.24
F	Т	Т	0.012
F	F	Т	0.108
F	Т	F	0.008
F	F	F	0.072

Table: Merging p(A, B, C)

ightharpoonup Marginalizing over variable B

Q #3b: Step-by-Step Derivation p(C|A)

Α	С	Prob
Т	Т	0.4
Т	F	0.4
F	Т	0.12
F	F	0.08

Table: Computing p(A, C)

ightharpoonup Marginalizing over variable B

Q #3b: Step-by-Step Derivation p(C|A)

С	(A)	Prob
Т	Т	0.4 / 0.8 = 0.5
F	Т	0.4 / 0.8 = 0.5
Т	F	0.12 / 0.2 = 0.6
F	F	0.08 / 0.2 = 0.4

Table: Computing p(C|A)

- ightharpoonup Marginalizing over variable B
- ightharpoonup Computing p(C|A)

Q #3b: Step-by-Step Derivation (Verification)

В	(A)	Prob
Т	Т	0.4
F	Т	0.6
Т	F	0.1
F	F	0.9

С	(A)	Prob
Т	T	0.5
F	T	0.5
Т	F	0.6
F	F	0.4

В	С	(A)	Prob
Т	Τ	Т	0.5 * 0.4 == 0.2
Τ	F	Т	0.5 * 0.4 == 0.2
F	Т	Τ	0.5 * 0.6 == 0.3
F	F	Т	0.5 * 0.6 == 0.3
Т	Т	F	0.6 * 0.1 == 0.06
Т	F	F	0.4 * 0.1 == 0.04
F	Т	F	0.6 * 0.9 == 0.54
F	F	F	0.4 * 0.9 == 0.36

All of the probabilities are equal, therefore ${\cal B}$ and ${\cal C}$ are independent given A.

CS 486/686: Intro to AI

Learning Goals

Unconditional and Conditional Independence

Examples of Bayesian Networks

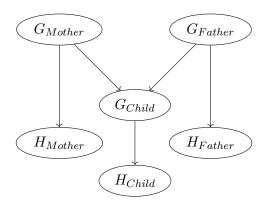
Why Bayesian Networks

Representing the Joint Distribution

Independence in Three Key Structures

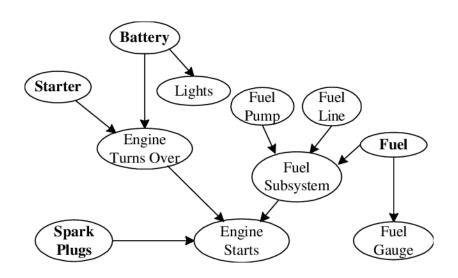
Revisiting Learning Goals

Inheritance of Handedness



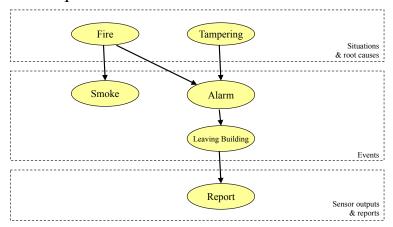
CS 486/686: Intro to AI Lecturer: Wenhu Chen 27 / 53

Car Diagnostic Network



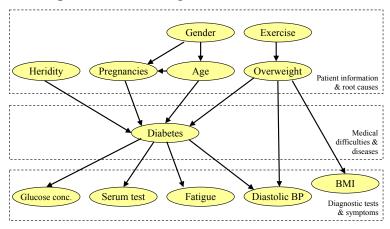
CS 486/686: Intro to AI Lecturer: Wenhu Chen 28 / 53

Example: Fire alarms



Report: "report of people leaving building because a fire alarm went off"

Example: Medical diagnosis of diabetes



Learning Goals

Unconditional and Conditional Independence

Examples of Bayesian Networks

Why Bayesian Networks

Representing the Joint Distribution

Independence in Three Key Structures

Revisiting Learning Goals

Why Bayesian Networks?

A probabilistic model of the Holmes scenario:

- ► The random variables: Earthquake, Radio, Burglary, Alarm, Watson, and Gibbon.
- # of probabilities in the joint distribution: $2^6 = 64$.
- For example,

$$\begin{split} &P(E \wedge R \wedge B \wedge A \wedge W \wedge G) = ? \\ &P(E \wedge R \wedge B \wedge A \wedge W \wedge \neg G) = ? \\ &\dots \text{ etc } \dots \end{split}$$

We can compute any probability using the joint distribution, but

- Quickly become intractable as the number of variables grows.
- Unnatural and tedious to specify all the probabilities.

Why Bayesian Networks?

A Bayesian Network

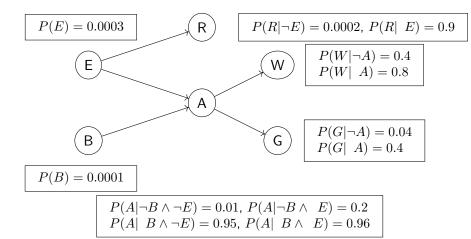
- is a compact version of the joint distribution
- ► takes advantage of the unconditional and conditional independence among the variables.

CS 486/686: Intro to Al Lecturer: Wenhu Chen 33 / 53

Reminder: Modelling the Holmes Scenario

- \rightarrow The random variables:
 - ► B: A Burglary is happening.
 - ► A: The alarm is going.
 - ▶ W: Dr. Watson is calling.
 - ► G: Mrs. Gibbon is calling.
 - ► E: Earthquake is happening.
 - R: A report of earthquake is on the radio news.

A Bayesian Network for the Holmes Scenario



How many probabilities do we need to encode the Network?

Bayesian Network

A Bayesian Network is a directed acyclic graph (DAG).

- Each node corresponds to a random variable.
- X is a parent of Y if there is an arrow from node X to node Y.
 - \rightarrow Like a family tree, there are parents, children, ancestors, descendants.
- ▶ Each node X_i has a conditional probability distribution $P(X_i|Parents(X_i))$ that quantifies the effect of the parents on the node.

Learning Goals

Unconditional and Conditional Independence

Examples of Bayesian Networks

Why Bayesian Networks

Representing the Joint Distribution

Independence in Three Key Structures

Revisiting Learning Goals

The Semantics of Bayesian Networks

Two ways to understand Bayesian Networks:

- ► A representation of the joint probability distribution
- ► An encoding of the conditional independence assumptions

CS 486/686: Intro to AI Lecturer: Wenhu Chen 38 / 53

The idea is that, given a random variable X, a small set of variables may exist that directly affect the variable's value in the sense that X is conditionally independent of other non-descendent variables given values for the directly affecting variables.

- Start with a set of random variables representing all the features of the model.
- ▶ Define the **parents** of random variable X_i , written as $parents(X_i)$.
- $ightharpoonup X_i$ is independent from other non-descendent variables given the $parents(X_i)$.

We can compute the full joint probability using the following formula, which is the product of multiple conditional probabilities.

$$P(X_n \wedge \cdots \wedge X_1) = \prod_{i=1}^n P(X_i | Parents(X_i))$$

CS 486/686: Intro to AI

Example: What is the probability that all of the following occur?

- The alarm has sounded
- Neither a burglary nor an earthquake has occurred
- Both Watson and Gibbon call and say they hear the alarm
- ► There is no radio report of an earthquake

Example: What is the probability that all of the following occur?

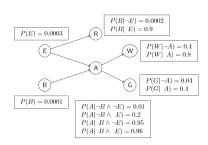
- The alarm has sounded
- Neither a burglary nor an earthquake has occurred
- Both Watson and Gibbon call and say they hear the alarm
- ► There is no radio report of an earthquake
- → Formulate as a joint probability:

$$P(\neg B \land \neg E \land A \land \neg R \land G \land W)$$
= $P(\neg B)P(\neg E)P(A|\neg B \land \neg E)P(\neg R|\neg E)P(G|A)P(W|A)$
= $(1 - 0.0001)(1 - 0.0003)(0.01)(1 - 0.0002)(0.4)(0.8)$
= 3.2×10^{-3}

Q #4: Calculating the joint probability

Q: What is the probability that all of the following occur?

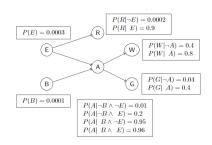
- ▶ NEITHER a burglary NOR an earthquake has occurred,
- ► The alarm has NOT sounded,
- ▶ NEITHER of Watson and Gibbon is calling, and
- ► There is NO radio report of an earthquake?
- (A) 0.5699
- (B) 0.6699
- (C) 0.7699
- (D) 0.8699
- (E) 0.9699



Q #4: Calculating the joint probability

Q: What is the probability that all of the following occur?

- NEITHER a burglary NOR an earthquake has occurred,
- ► The alarm has NOT sounded,
- NEITHER of Watson and Gibbon is calling, and
- There is NO radio report of an earthquake?
- (A) 0.5699
- (B) 0.6699
- (C) 0.7699
- (D) 0.8699
- (E) 0.9699



$$\rightarrow$$
 (A) $(1-0.0001)(1-0.0003)(1-0.01)(1-0.4)(1-0.04)(1-0.0002) = 0.5699$

CS 486/686: Intro to AI

Learning Goals

Unconditional and Conditional Independence

Examples of Bayesian Networks

Why Bayesian Networks

Representing the Joint Distribution

Independence in Three Key Structures

Revisiting Learning Goals

Burglary, Alarm and Watson



CS 486/686: Intro to AI Lecturer: Wenhu Chen 44 / 53

Q #5: Unconditional Independence

Q: Are Burglary and Watson independent?



- (A) Yes
- (B) No
- (C) Can't tell.

Q #5: Unconditional Independence

Q: Are Burglary and Watson independent?



- (A) Yes
- (B) No
- (C) Can't tell.

 \rightarrow Correct answer is *No*.

If you learned the value of B, would your belief of W change? If B is true, then Alarm is more likely to be true, and W is more likely to be true.

Q #6: Conditional Independence

Q: Are Burglary and Watson conditionally independent given Alarm?



- (A) Yes
- (B) No
- (C) Can't tell

Q #6: Conditional Independence

Q: Are Burglary and Watson conditionally independent given Alarm?



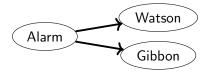
- (A) Yes
- (B) No
- (C) Can't tell
- \rightarrow Correct answer is *Yes*.

Assume that W does not observe B directly. W only observes A.

B and W could only influence each other through A.

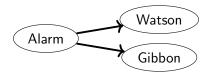
If A is known, then B and W do not affect each other.

Alarm, Watson and Gibbon



Q #7: Unconditional Independence

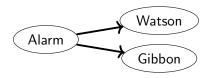
Q: Are Watson and Gibbon independent?



- (A) Yes
- (B) No
- (C) Can't tell

Q #7: Unconditional Independence

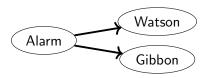
Q: Are Watson and Gibbon independent?



- (A) Yes
- (B) No
- (C) Can't tell
- ightarrow Correct answer is No. If Watson is more likely to call, then Alarm is more likely to go off, which means that Gibbon is more likely to call.

Q #8 Conditional Independence

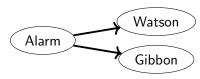
Q: Are Watson and Gibbon conditionally independent given Alarm?



- (A) Yes
- (B) No
- (C) Can't tell

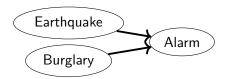
Q #8 Conditional Independence

Q: Are Watson and Gibbon conditionally independent given Alarm?



- (A) Yes
- (B) No
- (C) Can't tell
- ightarrow Correct answer is Yes. Watson and Gibbon are both unreliable sensors for Alarm. If Alarm is known, then Watson and Gibbon do not affect each other.

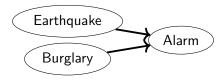
Earthquake, Burglary, and Alarm



CS 486/686: Intro to AI Lecturer: Wenhu Chen 50 / 53

Q #9 Unconditional Independence

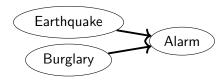
Q: Are Earthquake and Burglary independent?



- (A) Yes
- (B) No
- (C) Can't tell

Q #9 Unconditional Independence

Q: Are Earthquake and Burglary independent?

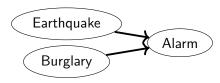


- (A) Yes
- (B) No
- (C) Can't tell

 \rightarrow Correct answer is Yes.

Q #10: Conditional Independence

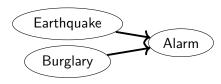
Q: Are Earthquake and Burglary conditionally independent given Alarm?



- (A) Yes
- (B) No
- (C) Can't tell

Q #10: Conditional Independence

Q: Are Earthquake and Burglary conditionally independent given Alarm?



- (A) Yes
- (B) No
- (C) Can't tell

 \rightarrow Correct answer is No. Suppose that the Alarm is going. If there is an Earthquake, then it is less likely that the Alarm is caused by Burglary. If there is a Burglary, it is less likely that the Alarm is caused by Earthquake.

Revisiting Learning Goals

- Given a probabilistic model, determine if two variables are unconditionally independent, or conditionally independent given a third variable.
- Give examples of deriving a compact representation of a joint distribution by using independence assumptions.
- Describe components of a Bayesian network.
- Compute a joint probability given a Bayesian network.
- Explain the independence relationships in the three key structures.

CS 486/686: Intro to AI Lecturer: Wenhu Chen 53 / 53